

CS 6320 — Assignment 5

Name: Brandon Jackson

Semester: Summer 2026

Course: Deep Learning (CS 6320)

Part A — Model evaluation and generalization evidence (BGG portfolio)

Work type: Portfolio baseline evaluation (not approved practice work).

Task and model summary

Item	Lock
Dataset	Kaggle <i>Board Games Database from BoardGameGeek</i> (<code>games.csv</code> , 21,925 rated games)
Unit	One row per game (BGGId)
Target	<code>high_rating = 1</code> if <code>AvgRating >= 7.0</code> , else <code>0</code> (~26.9% positive)
Stakeholder use	Pre-release / preorder stocking screen for a hobby retailer
Model evaluated	Logistic regression (<code>SGDClassifier</code> , log-loss, L2) — Assignment 4 baseline stage
Primary intervention	Balanced class weights via <code>sample_weight</code> (addresses minority-class recall)
Seed	6320
Implementation	6320-hw5/scripts/; end-to-end: <code>bash run_local.sh</code>

Features follow the locked Assignment 4 charter: numeric design metadata, `GameWeight`, category flags, simple text-length features, parsed `goodplayers_count`. Hard-excluded rating fields, rank columns, popularity counts, and mis-timed demand fields (`NumOwned`, `NumWant`, `NumWish`).

Split strategy and audit evidence

Strategy: stratified random 70% / 15% / 15% hold-out by game row (each row is one game — no duplicate entities across splits). Seed 6320. Model selection uses **validation**; **test** is held out for a final check only.

Why appropriate: Matches the charter’s hold-out-by-game plan for a static snapshot where each game appears once. Stratification preserves ~27% positive rate in each split.

Limitation: Random stratified split does **not** simulate the charter’s pre-release prediction moment; a **time-based** split by YearPublished remains a Week 5+ follow-up.

Split	Rows	% of data	Positive rate
train	15,347	70.0%	26.9%
validation	3,289	15.0%	27.7%
test	3,289	15.0%	26.2%

Plausibility check: positive rates are stable across splits (within ~1.5 pp), so the stratified partition did not collapse class balance. Numeric metadata distributions by split are saved in prep/bgg_split/split_audit/split_numeric_distributions.csv (e.g. YearPublished, GameWeight means are comparable across train/validation/test).

Run evidence: python3 scripts/prepare_bgg_data.py; python3 scripts/run_split_audit.py; artifacts under prep/bgg_split/.


Aggregate metrics

Majority-class baseline (always predict not-high): validation accuracy **72.3%**, F1 **0.0**, recall **0.0** on high_rating=1.


Run	Split	Accuracy	Precision (pos)	Recall (pos)	F1 (pos)	ROC-AUC
Baseline (unweighted)	validation	0.790	0.687	0.442	0.538	0.814
Baseline (unweighted)	test	0.803	0.678	0.473	0.557	0.831
Intervention (balanced weights)	validation	0.748	0.534	0.714	0.611	0.815
Intervention (balanced weights)	test	0.762	0.532	0.753	0.624	0.832

The baseline beats the majority class on F1 and ROC-AUC but **under-recalls** high-rated games. The intervention trades some accuracy/precision for much better minority recall while holding ROC-AUC flat.


Learning curves (train vs validation)




Training loss



Validation loss



Validation F1



Validation recall

Figures 1–4: Baseline vs balanced-weight intervention; 40 SGD epochs, learning rate 0.001.

Generalization diagnosis: Train and validation **loss/F1 track closely** for both runs (no large train–val divergence). Evidence suggests **plausible generalization at this capacity**, not severe overfitting. The main failure mode is **class imbalance handling**, not runaway train fit: baseline validation recall stays ~0.44 while accuracy looks acceptable.

Generalization intervention

Report field	Content
Observed issue	Unweighted baseline favors the majority class — high accuracy (79%) but low recall (44%) on high_rating=1, poor retailer utility for “don’t miss hits.”

Report field	Content
Expected test	Balanced class weights should raise recall/F1 on the minority class without requiring a new architecture.
Intervention	Same features and L2 ($\alpha=1e-4$); apply <code>compute_class_weight('balanced')</code> through <code>sample_weight</code> on training only.
Before (validation)	F1 0.538 , recall 0.442 , precision 0.687
After (validation)	F1 0.611 , recall 0.714 , precision 0.534
What changed	Better minority detection; lower precision and accuracy — expected trade-off for imbalance. ROC-AUC unchanged (~0.815).
Limitation	Single linear baseline; no time-split; calibration still poor (below).

Error and slice analysis (intervention model, validation)

Confusion matrix (validation):

	Pred 0	Pred 1
True 0	1,811	568
True 1	260	650

Slice — GameWeight (complexity):

Bin	n	Positive rate	F1 (high_rating)	Recall (high_rating)
light (<2)	1,698	11.8%	0.31	0.28
medium	1,134	36.2%	0.58	0.73
heavy	392	63.5%	0.78	0.98
very heavy	65	78.5%	0.88	1.00

Interpretation: The model works best on heavier games where the positive class is more common; **light games** remain hard — consistent with weak metadata signal and representativeness limits in the charter.

Year-level slices (`outputs/error_analysis/slice_by_year.csv`) show small-n bins; several older/low-count years have unstable precision. High-confidence errors (confidence ≥ 0.8 and wrong): saved in `outputs/error_analysis/high_confidence_errors.csv`.

Calibration

Calibration is **meaningful** for this binary screening task (probabilities drive stock/skip thresholds).

Decile reliability bins (outputs/error_analysis/calibration_bins.csv) show **systematic overconfidence**: predicted probabilities exceed observed positive rates in most bins (e.g. mean predicted **0.52** vs observed **0.30** in the 0.47–0.58 bin; gap up to ~0.22 in mid bins). The balanced-weight model improves recall but **does not yet produce trustworthy probability estimates** for deployment thresholds.

Run evidence (local)

- Command: `bash run_local.sh` from repo root
- Data: 6320-hw2/part_b/data/bgg/games.csv (or GAMES_CSV)
- Artifacts: prep/bgg_split/manifest.json, outputs/*/history.csv, outputs/plots/*.png, outputs/error_analysis/*.csv
- Compute: local CPU only (~seconds); CHPC not required for this dataset/model size

Part B — Portfolio audit update and reliability judgment

Audit risk trace (Assignment 4 → Week 5 evidence)

A4 risk / success criterion	Week 5 evidence	Status
Class imbalance (~27% positive) — report F1/recall, not accuracy alone	Majority F1=0; baseline acc 79% but recall 44%; intervention recall 71% val / 75% test	Confirmed — imbalance distorts accuracy; weighting helps recall
Leakage / feature timing — exclude rating & mis-timed demand fields	Manifest lists exclusions; model uses design metadata only	Partially tested — exclusions enforced in prep; time-stamp audit not yet done
Representativeness — BGG enthusiast bias; not universal quality	Slices show weak performance on light games and sparse year bins	Confirmed risk in practice
Success: classical baseline beats majority	ROC-AUC ~0.81 vs 0.5; F1 > 0 vs 0	Met
Honest hold-out evaluation	Stratified split audit; test metrics reported once	Met for random split; time-split still open
Responsible use — screening aid, human review	Calibration gaps + slice weakness argue against automated decisions	Reinforced

Reliability judgment (bounded, current stage)

Supported now: A leakage-aware **tabular baseline** can rank games better than chance (ROC-AUC ~0.83 test) and better than always predicting “not high.” Balanced weighting improves **minority recall** for a “don’t miss hits” retailer story. Train/validation curves do not show runaway overfitting at this model size.

Not supported: Deployment-ready stocking decisions. Probability outputs are **miscalibrated**; **light-weight** games are poorly served; random split does not prove pre-release generalization; no mechanics join beyond category flags yet; no human-in-the-loop workflow tested.

Recommendation: Use for internal screening only — flag titles for human review, not auto stock/skip. Prefer intervention weights when recall matters more than precision; do not trust raw probability thresholds until recalibration.

Evidence still needed before client-facing use

- Time-based split / recent-release cohort evaluation (charter follow-up)
- Feature-timing audit on any newly added fields
- Recalibration (Platt scaling or isotonic) on validation, checked on test
- Gradient-boosted trees (A4 initial candidate) vs logistic — next staged model
- Error review with publisher/franchise grouping to catch memorization via title text proxies

Next staged model improvement (from A4 plan)

Stage	Next step	Why
Baseline	Done — majority + logistic with imbalance diagnosis	Establishes floor and metric story
Initial candidate	Gradient-boosted trees on same feature manifest	Capture nonlinear metadata/category interactions
Revised alternative	Recalibrated probabilities or simpler logistic if trees overfit	Address calibration gap
Final evidence	Time-split + stakeholder cost matrix (false stock vs false skip)	Tie metrics to retailer decision

AI disclosure

Tool: Cursor (Composer)

AI-assisted: Repo scaffolding, sklearn training/evaluation scripts, Week 5 example patterns, first-draft writeup prose.

My work: Locked portfolio path (no TLC practice); chose balanced-weight intervention from baseline recall failure; ran and verified all local metrics, slices, and calibration tables against saved outputs.

Certification: I certify that all work not described above is my own, and I verified AI-generated content for accuracy.